

# Managing Computationally Expensive Blackbox Multiobjective Optimization Problems with libEnsemble

Tyler Chang<sup>1</sup>, Jeffrey Larson<sup>2</sup>, Layne Watson<sup>1</sup>, & Thomas Lux<sup>1</sup>

<sup>1</sup> Dept. of Computer Science, Virginia Polytechnic Institute & State Univ.

<sup>2</sup> Mathematics and Computer Science Division, Argonne National Laboratory

May 19, 2020



Background in MOPs



The Algorithm

Parallel implementations

Results

- ▶ The Multiobjective Optimization Problem (MOP) generalizes the Single Objective (Scalar) Optimization Problem (SOP);
- ▶ The MOP attempts to balance the tradeoff between multiple conflicting objectives;
- ▶ Whereas the SOP generally has a unique solution, the solution to a MOP is a *set of Pareto optimal* solutions;

- ▶ The solution to a MOP is a set of *nondominated* or *Pareto optimal* solutions;
- ▶  $x^*$  is Pareto optimal if for all  $x \in \mathcal{X}$ ,  $F(x) \not\leq F(x^*)$ ;



Find a discrete set of approximately nondominated objective points that describes the Pareto front, and the corresponding efficient designs

## Types of MOPs

functions are “cheap” to evaluate derivative info is available	functions are “cheap” to evaluate no derivative info is available
functions are costly to evaluate derivative info is available	functions are costly to evaluate no derivative info is available

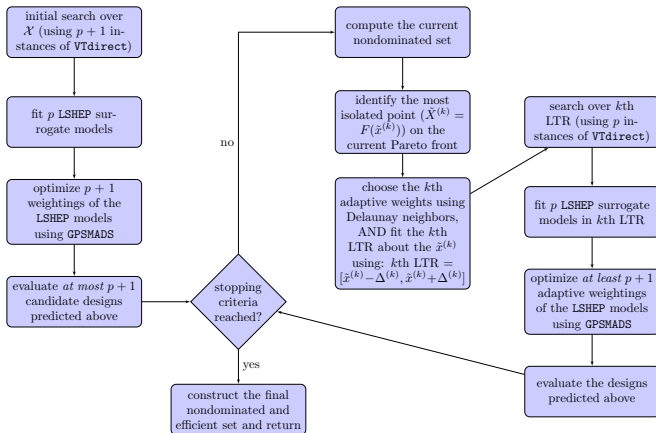
Focus on bottom right: **expensive blackbox MOPs!**

Provide a parallel implementation of the multiobjective optimization algorithm (MOA) proposed by

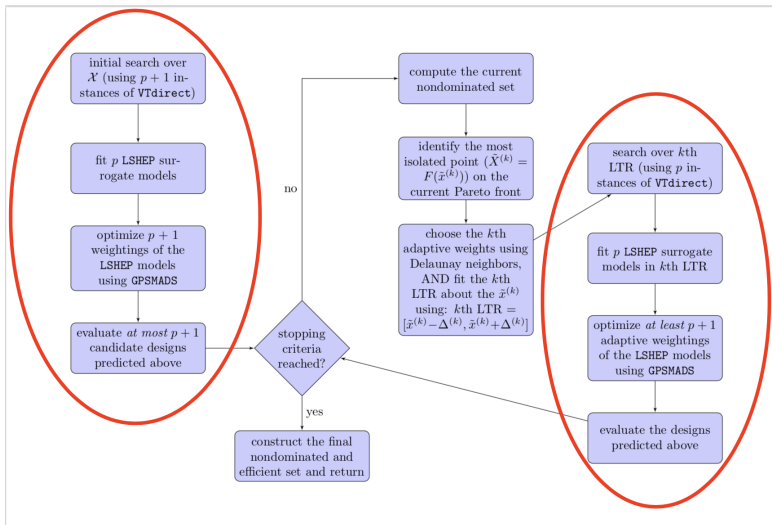
*Shubhangi Deshpande, et al. "Multiobjective optimization using an adaptive weighting scheme." Optimization Methods and Software 31.1 (2016): 110-133.*

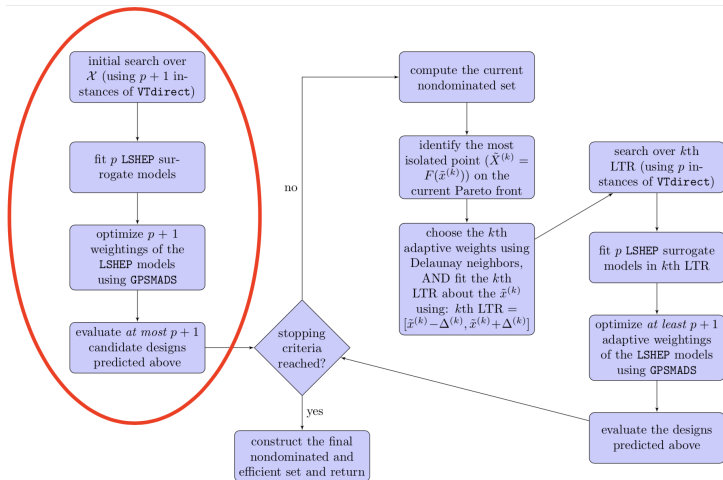
Combines adaptive weighting scheme, response surface modeling, and trust region methods

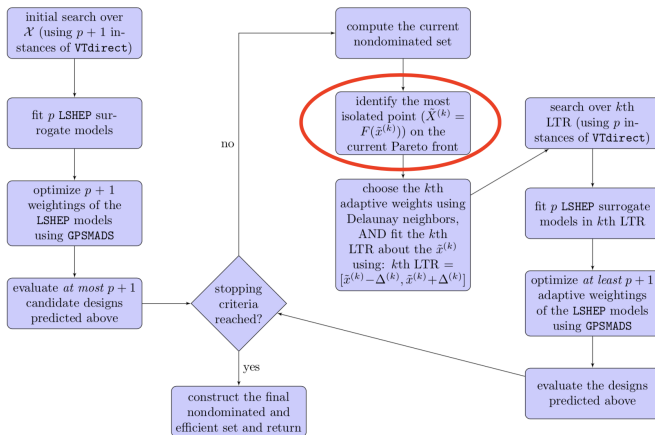
# The Algorithm Outline



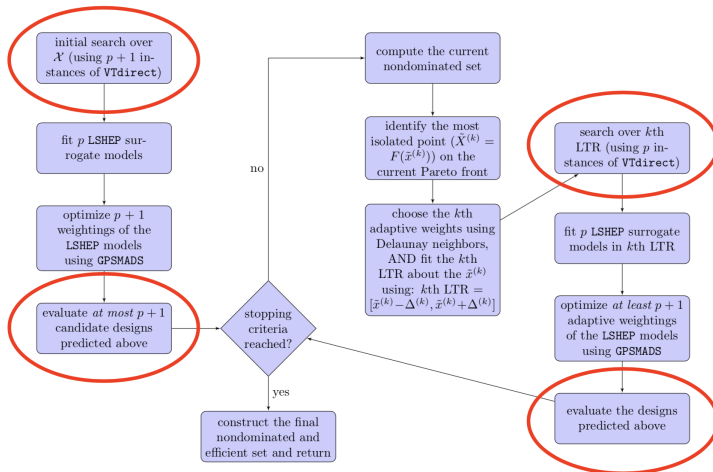








1. The function  $F$  (left to the user)
2. Iteration complexity (assumed to not offer much improvement)
3. Function evaluations  $\leftarrow$  **most important**



- ▶ Recall that  $F$  is being distributed by user
- ▶ Use OpenMP *shared memory parallelism*, essentially for achieving asynchronous behavior
- ▶ Puts burden of distribution on user
- ▶ Better flexibility for real-world HPC systems

- ▶ Recall that  $F$  is being distributed by user
- ▶ Use OpenMP *shared memory parallelism*, essentially for achieving asynchronous behavior
- ▶ Puts burden of distribution on user
- ▶ Better flexibility for real-world HPC systems

To parallelize function evaluations:

- ▶ Batch of candidates can be trivially evaluated in parallel
- ▶ Instead of using pVTdirect from VTDIRECT95 (which uses MPI) made a modification to the serial code called bVTdirect that evaluates batches of points in parallel using OpenMP
- ▶ For maximal parallelism, replace bVTdirect with a Latin hypercube design, which can be evaluated in parallel

The `libEnsemble` library is part of the Exascale Computing Project at Argonne to harness increased levels of concurrency when distributing large simulations over extreme scale resources

- ▶ Generator function (generates simulations to run)
- ▶ Simulator function (runs simulations, possibly in parallel)
- ▶ Allocator function (decides whether to generate or simulate)



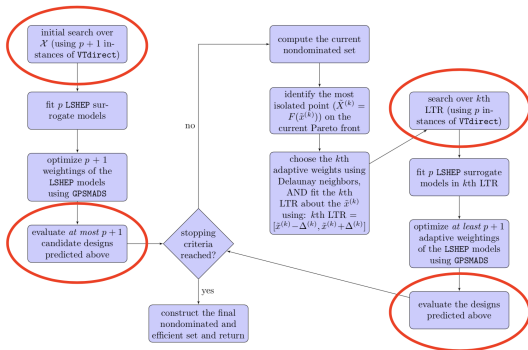
The `libEnsemble` library is part of the Exascale Computing Project at Argonne to harness increased levels of concurrency when distributing large simulations over extreme scale resources

- ▶ Generator function (generates simulations to run)
- ▶ Simulator function (runs simulations, possibly in parallel)
- ▶ Allocator function (decides whether to generate or simulate)

VTMOP is the *generator* for `libEnsemble`


- ▶ Each call to the *generator* runs a half-iteration and requests a design exploration (using Latin hypercube) or a batch of candidate evaluations
- ▶ The *simulator* evaluates all the requested designs
- ▶ The *allocator* waits until all designs are evaluated and then swithes back to the *generator*

- ▶ Want nice big batches that match available resources
- ▶ Use a Latin hypercube search during the search phase
- ▶ Pad out batches of candidates using additional weights




- ▶  $F^{(c)}(x) = (\|x - 0.5e^{(1)}\|_2^2, \dots, \|x - 0.5e^{(p)}\|_2^2)$
- ▶ Convex Pareto front  $\Rightarrow$  “easier” problem

Tradeoff curve between objectives F1, F2, and F3



- ▶ DTLZ2 from Deb et al.
- ▶ Concave Pareto front  $\Rightarrow$  “harder” problem

Tradeoff curve between objectives F1, F2, and F3



Evaluating a multiobjective optimization solution is a multiobjective problem...

We want:

1. Many points on the Pareto front (measured by *cardinality* of the solution set)
2. Good convergence of the solution points to the true Pareto front (measured by RMSE)
3. Even spacing/good coverage of the solution set (measured by Delaunay discrepancy, computed with SciPy)

Average (5 runs) number of solutions, RMSE, and Delaunay discrepancy for  $F^{(c)}$  and DTLZ2 with  $d = 5$  for VTMOOP using bVTdirect (DIR), Latin hypercube search (LH), and libEnsemble (LIBE)


Prob/Meth	$p = 2$	$p = 3$	$p = 4$
$F^c$ / bVTdir	73, .00100, .207	173, .0505, .579	288, .101, NA
$F^c$ / libE1	38, .0115, .180	93, .0665, .512	171, .117, .689
$F^c$ / libE2	78, .0127, .158	189, .0560, .429	283, .104, .551
DTLZ2 / bVTdir	139, .00713, .109	354, .0401, .230	658, .0443, NA
DTLZ2 / libE1	80, .0993, .139	264, .167, .458	510, .210, .757
DTLZ2 / libE2	66, .103, .201	258, .175, .691	548, .201, .793

NA = Missing values due to SciPy's Delaunay triangulation tool


CPU time / wall time in seconds for DIR, LH, and LIBE versions with 36 cores & 1 sec (+ var)/1 core per eval

$p$	Meth	$F_c$ no var	$F_c$ + var	DTLZ2 no var	DTLZ2 + var
2	bVTdir1	2008 / 1037	2007 / 1039	2007 / 1093	2004 / 1082
	bVTdir2	NA / 170	NA / 239	NA / 175	NA / 240
	libE1	2015 / 89	2017 / 107	2028 / 89	2011 / 109
	libE2	2051 / 112	2070 / 142	2060 / 111	2064 / 143
3	bVTdir1	2012 / 717	2012 / 719	2021 / 797	2018 / 797
	bVTdir2	NA / 137	NA / 207	NA / 165	NA / 237
	libE1	2023 / 94	2034 / 116	2040 / 95	2023 / 117
	libE2	2077 / 133	2066 / 144	2054 / 99	2057 / 126
4	bVTdir1	2026 / 582	2029 / 586	2177 / 807	2149 / 782
	bVTdir2	NA / 134	NA / 208	NA / 280	NA / 348
	libE1	2042 / 108	2045 / 127	2176 / 200	2201 / 280
	libE2	2134 / 190	2124 / 186	2182 / 227	2185 / 257


# CPU Usage (36 cores)



bVTdir1 (not suitable for 36 cores)



libE1



libE2

Background in MOPs

The Algorithm

Parallel implementations

Results