

ParMOO: A Python library for parallel multiobjective simulation optimization

Tyler Chang^a and Stefan Wild^{a,b}

^aMathematics and Computer Science Division,
Argonne National Laboratory

^bApplied Mathematics and Computational Research Division,
Lawrence Berkeley National Laboratory

SIAM CSE 23

Outlines

Introduction to MOSO + my experience

3 challenges + solutions

ParMOO software design + release

Example Problems

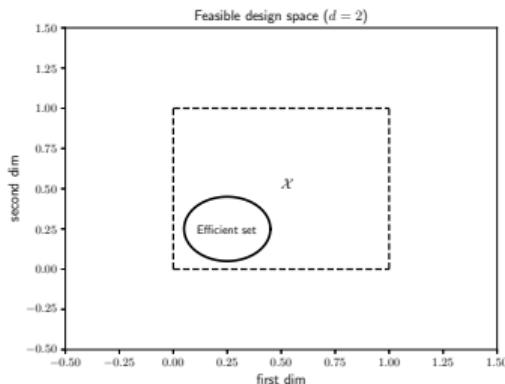
Conclusions + some closing thoughts

Multiobjective Optimization Problems

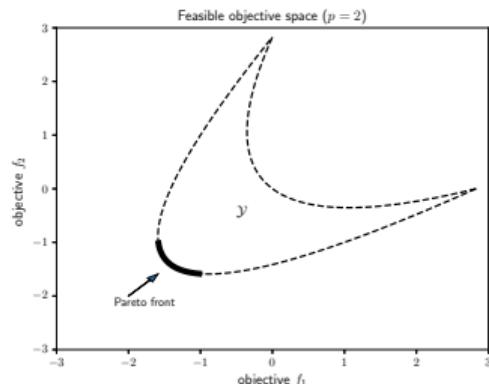
$$\min_{x \in \mathcal{X}} F(x)$$

Multiobjective Optimization Problems

$$\min_{x \in \mathcal{X}} F(x)$$

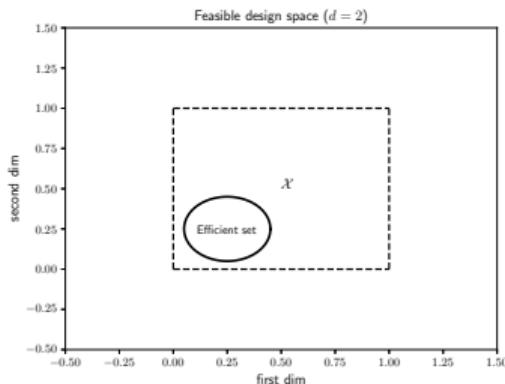


$$F : \mathcal{X} \rightarrow \mathcal{Y}$$



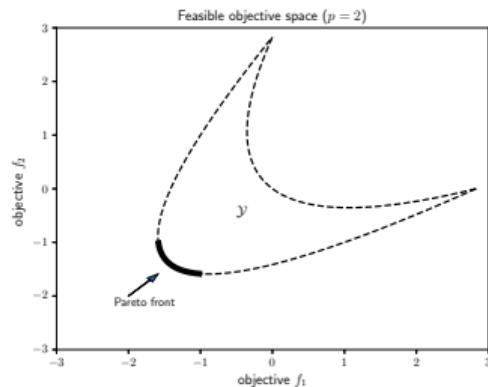
Multiobjective Optimization Problems

$$\min_{x \in \mathcal{X}} F(x)$$



$$F : \mathcal{X} \rightarrow \mathcal{Y}$$

expensive
blackbox process

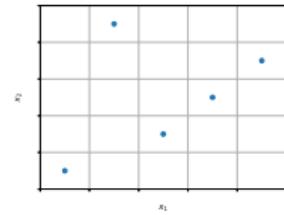


Multiobjective Response Surface Methodology

or Model-Based Optimization or Active Learning

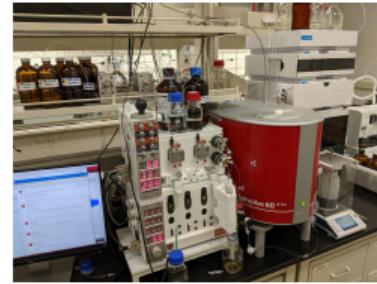
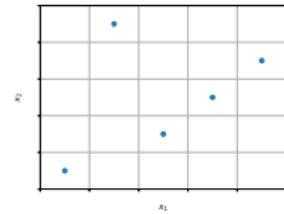
Multiobjective Response Surface Methodology

or Model-Based Optimization or Active Learning



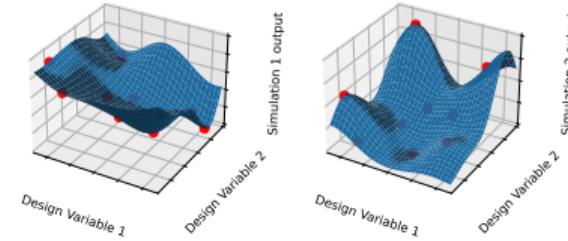
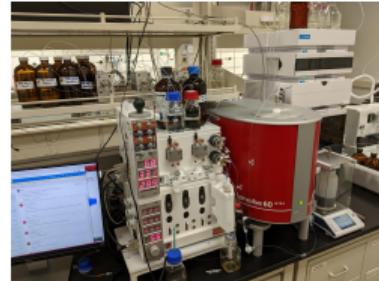
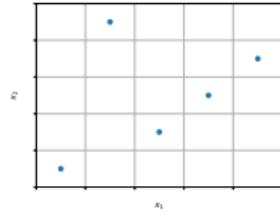
Multiobjective Response Surface Methodology

or Model-Based Optimization or Active Learning



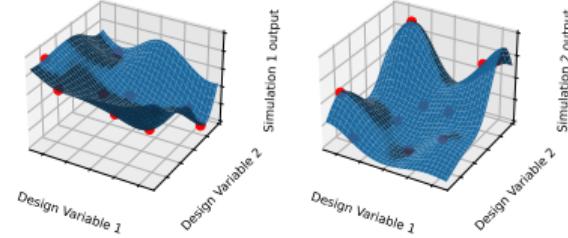
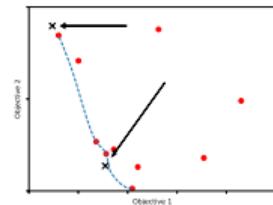
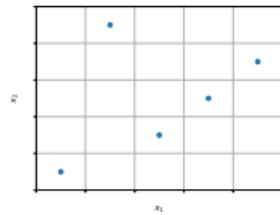
Multiobjective Response Surface Methodology

or Model-Based Optimization or Active Learning



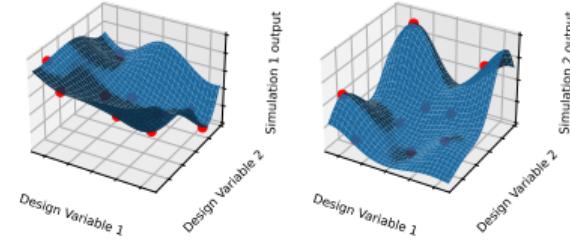
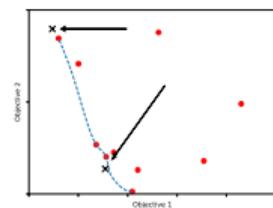
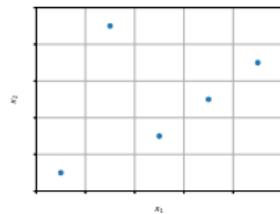
Multiobjective Response Surface Methodology

or Model-Based Optimization or Active Learning

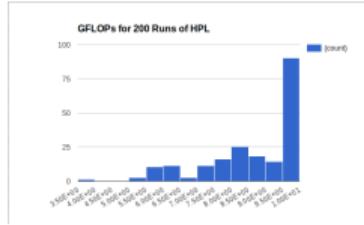


Multiobjective Response Surface Methodology

or Model-Based Optimization or Active Learning



Example: HPC Performance Tuning

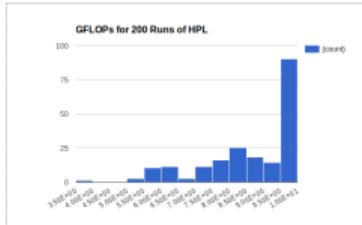


VT VarSys Project – 40 runs of HPL



ANL – LCRC Computing Resources: Bebop

Example: HPC Performance Tuning



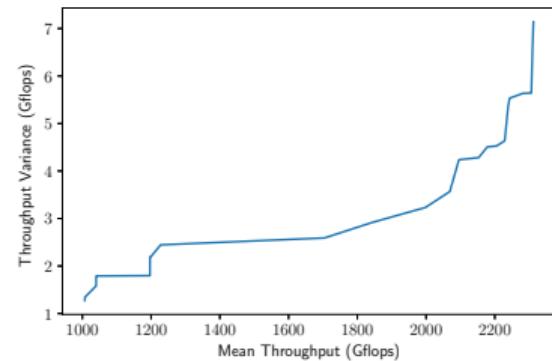
VT VarSys Project – 40 runs of HPL



ANL – LCRC Computing Resources: Bebop

```
HPL.dat
HPL (High Performance Linpack) benchmark input file
Copyright (c) 2002, University of Tennessee
HPL.out    output file name (if any)
0          generate output (1=yes,0=no)
0          number of problems solved (0)
29 30 34 35  Ns
0          # of RHS
1 2 3 4  NBs
0          # of process mapping (NbRows,1=column-major)
3          # of process grids (P > 0)
2 4 4  Ps
0          threshold
3          # of nested fact
0 3 2  PRFACTs (0:left, 3:center, 2:right)
2          # of recursive stopping criterion
NMIN
0          # of panels in recursion
NMAX
0          # of recursive panel fact
NFACTS (0:left, 3:center, 2:right)
1 2  # of broadcast
0          # of broadcast
1 2  RCOND (1=left,1=center,2=right,3=2TM,4=LSq,5=QR)
1 0  DGTols (>=0)
0 0  DGEQRF (1=col,1=diag,2=col)
0 0  LS (1=transposed,2=non-transposed,3=non-transposed,4=non-transposed)
0 0  U (1=triangular,2=non-triangular,3=non-triangular)
1 0  Equilibrate (1=none,2=sym)
0 0  memory alignment (in double C > 0)
```

VTMOP solver



[1] Chang, Larson, and Watson. Multiobjective optimization of the variability of the high-performance LINPACK solver. In Proc. WSC 2020.

Challenge 1:

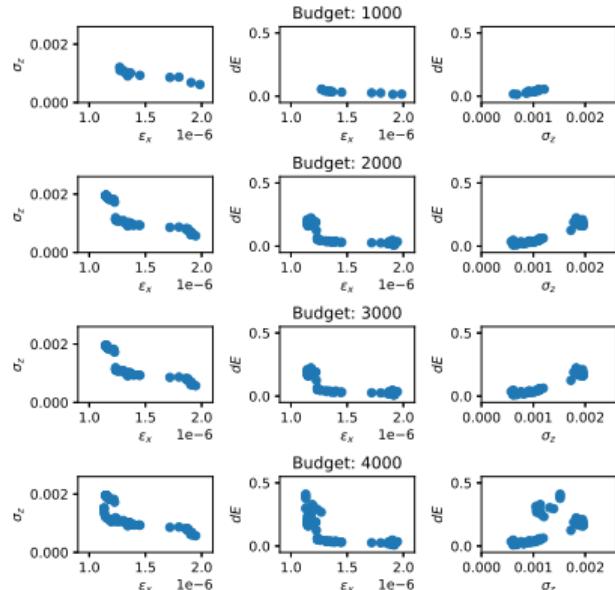
mixed vars + problem types

Example: Particle Accelerator Design

OPAL-t
(Object Oriented Parallel
Accelerator Library for
beam-lines + linacs)

```
OPAL.dat
MULinac benchmark input file
Developed by the Accelerator Laboratory, University of Tennessee
10.1.0.0          default float type (double)
0               device set (bsl0000,bsl0001,filg)
4               # of problems sizes (M)
20 30 35      Ns
2               # of Ns
1 2 3 4        Ns
0               # of MPI
3               # of process mapping (bottom-left=column-major)
2 1 4           Ps
0               Qs
10.0
threshold
1               # of parallel fact
0 3 2           PMFACT (0=left, 3=right, 2=right)
2 4           # of recursive stopping criterium
NODIV
1               # of panels in recursion
NODIV
3 2           # of recursive panel fact.
PMFACT (0=left, 3=right, 2=right)
1               # of parallel fact
BCM3D (0=1g_1,1=1g_1+2g_2,2=2g_1,3=2g_2,4=2g_3,5=2g_4)
1               # of parallel fact
DEPTHS (0=0)
2               # of parallel fact
2mix
Swapping (0=0,1=1,2=2,3=3,4=4)
L1 in (0=transposed,1=non-transposed) Term
0               L2 in (0=transposed,1=non-transposed) Term
Equilibrium (0=none,1=real) Term
memory alignment in double (> 0)
```

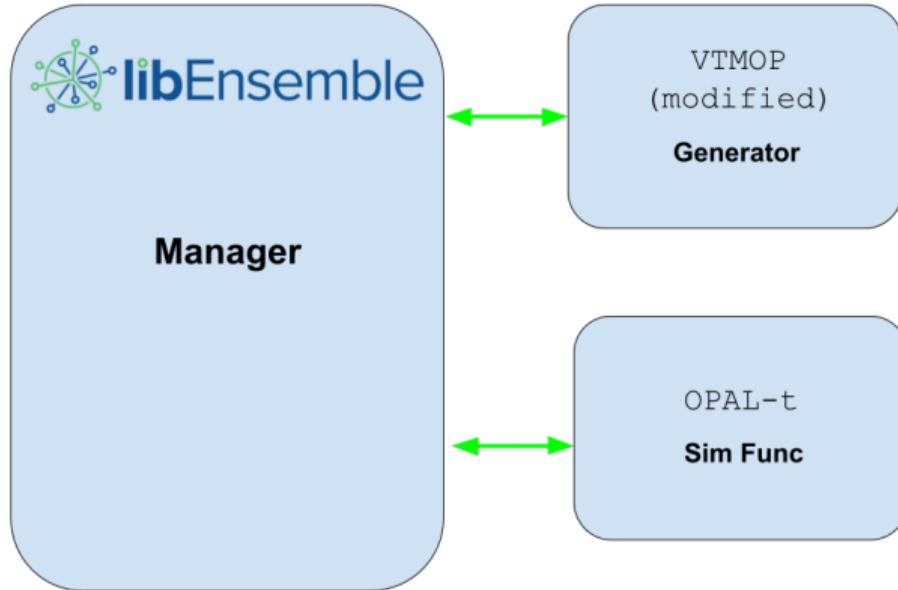
VTMOP solver



[2] Chang et al. Algorithm 1028: VTMOP: Solver for blackbox multiobjective optimization problems. ACM TOMS 48(3):36 (2022).

[3] Neveu et al. Comparison of multiobjective optimization methods for the LCLS-II photoinjector. CPC 283:108566 (2023).

Handling parallel evals



[4] Chang et al. Managing computationally expensive blackbox multiobjective optimization problems with libEnsemble. In Proc. SpringSim 2020.

Challenge 2: **parallel evals + computing environments**

Commercial solutions

Commercial solutions

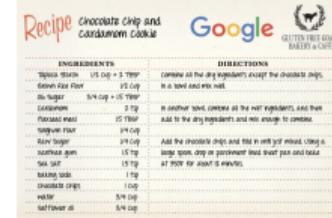


"Using Bayesian optimization for balancing metrics in recommendation systems" by Yunbao Ouyang et al. on LinkedIn Engineering Blog.

Commercial solutions



"Using Bayesian optimization for balancing metrics in recommendation systems" by Yunbao Ouyang et al. on LinkedIn Engineering Blog.

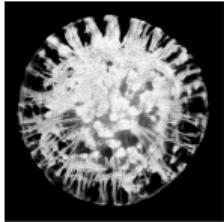


"The makings of a smart cookie" by Daniel Golovin on Google Research Blog.

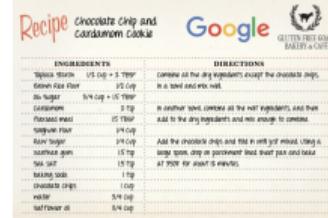
Commercial solutions



"Using Bayesian optimization for balancing metrics in recommendation systems" by Yunbao Ouyang et al. on LinkedIn Engineering Blog.



"Accelerating molecular optimization with AI" by Payel Das et al. on IBM Research Blog.

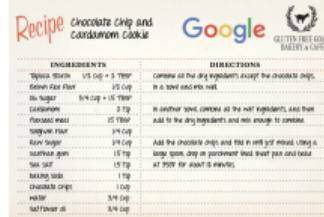


"The makings of a smart cookie" by Daniel Golovin on Google Research Blog.

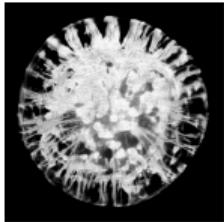
Commercial solutions



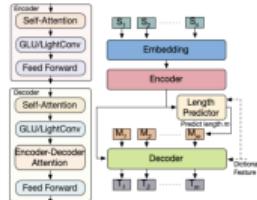
"Using Bayesian optimization for balancing metrics in recommendation systems" by Yunbao Ouyang et al. on LinkedIn Engineering Blog.



"The makings of a smart cookie" by Daniel Golovin on Google Research Blog.



"Accelerating molecular optimization with AI" by Payel Das et al. on IBM Research Blog.

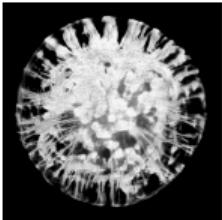


"Optimizing model accuracy and latency using Bayesian multi-objective NAS" by David Eriksson et al. on Meta Research Blog.

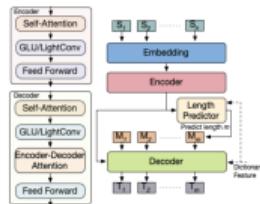
Commercial solutions



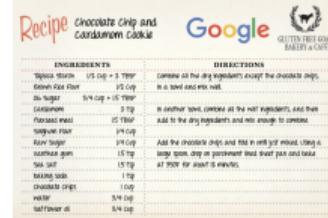
"Using Bayesian optimization for balancing metrics in recommendation systems" by Yunbao Ouyang et al. on LinkedIn Engineering Blog.



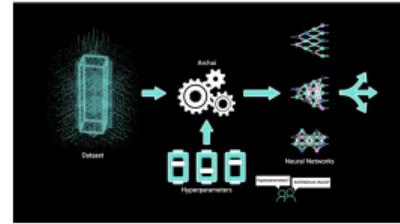
"Accelerating molecular optimization with AI" by Payel Das et al. on IBM Research Blog.



"Optimizing model accuracy and latency using Bayesian multi-objective NAS" by David Eriksson et al. on Meta Research Blog.



"The makings of a smart cookie" by Daniel Golovin on Google Research Blog.



"Archai can design your neural network with state-of-the-art NAS" by Shital Shah et al. on Microsoft Research Blog.

SOA and structure-exploiting blackbox optimization

SOA and structure-exploiting blackbox optimization



"Optimization and root finding (scipy.optimize)" in SciPy v1.10.0 API.

SOA and structure-exploiting blackbox optimization



"Optimization and root finding (`scipy.optimize`)" in SciPy v1.10.0 API.



Stochastic dimension reduction explained in this context by Stefan [5].

[5] Wild. Optimization and learning with zeroth-order stochastic oracles. *SIAM News* 56(1):1,3 (2023).

SOA and structure-exploiting blackbox optimization



"Optimization and root finding (`scipy.optimize`)" in SciPy v1.10.0 API.



Stochastic dimension reduction explained in this context by Stefan [5].



SOS structure can be exploited by DFO solver POUNDERS in TAO [6].

[5] Wild. Optimization and learning with zeroth-order stochastic oracles. *SIAM News* 56(1):1,3 (2023).

[6] Wild. Solving derivative-free nonlinear least squares problems with POUNDERS. *Adv. and Trends in Optimization with Engineering Applications*, Ch. 40, pp. 529–540 (2017).

Challenge 3:

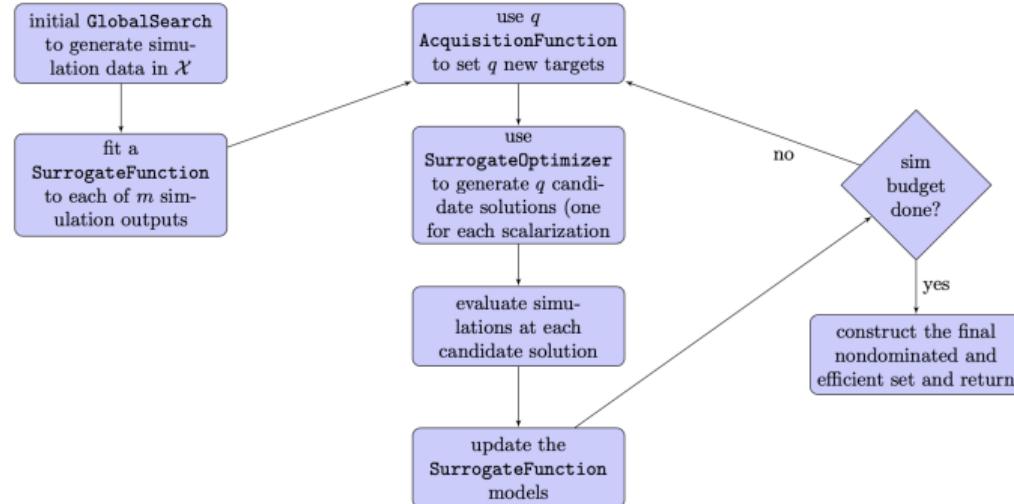
SOA optimization + exploiting problem structure

Design goals:

1. Highly customizable framework for multiobjective RSM
2. Flexible problem types (mixed-variables, constraints, etc.)
3. Easy to use, deploy, and extend (unforeseen use-cases and environments)
4. Exploit structure and domain knowledge

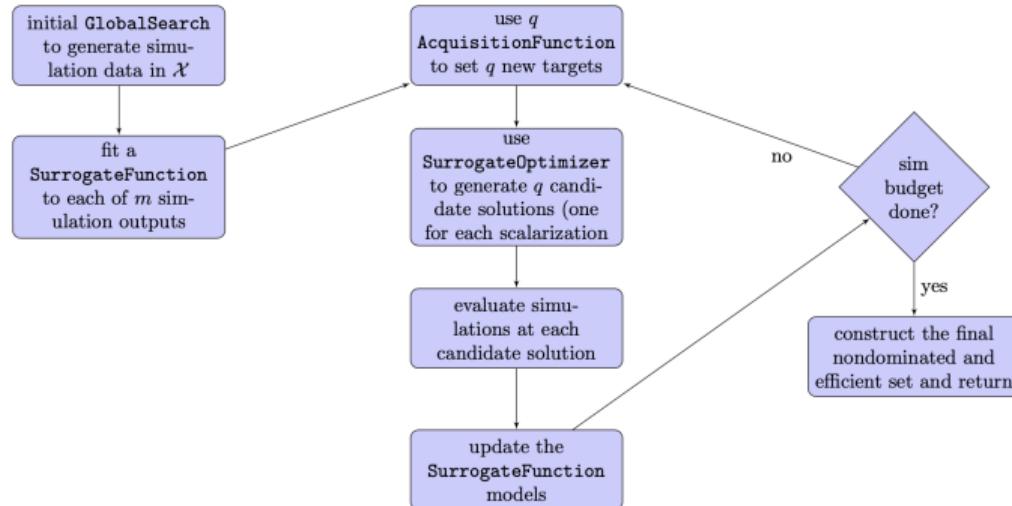
Customizability

ParMOO uses an object-oriented framework:



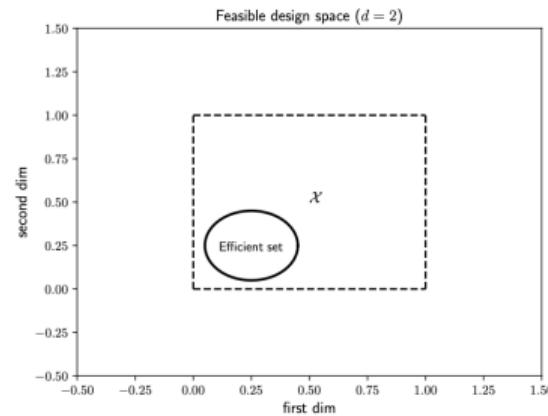
Customizability

ParMOO uses an object-oriented framework:



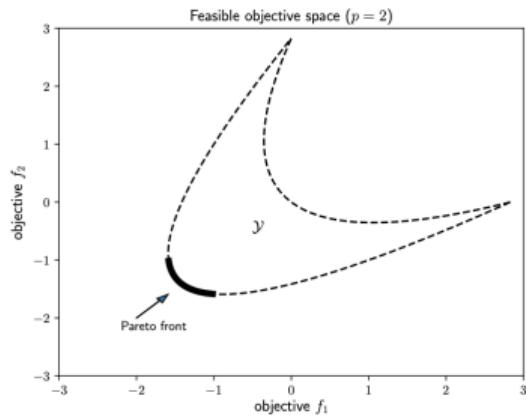
- ▶ Search/DOE
- ▶ Surrogate model
- ▶ Acquisition function
- ▶ Single-obj solver

Simulation Structure



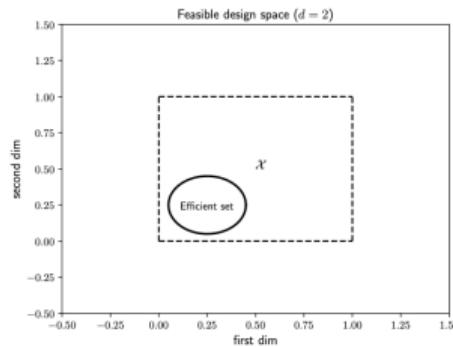
Design space

Objective Functions
→



Objective space

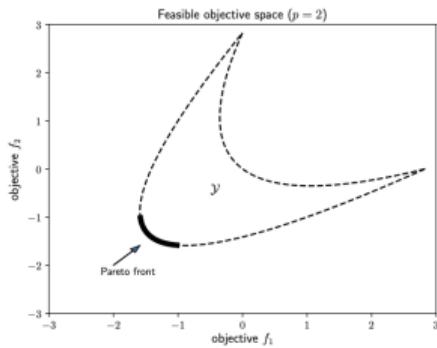
Simulation Structure



Design space

Simulations → \mathcal{S}

Objectives →



Objective space

Simulation Structure

Simulation Structure

$$f_i(x) = h_i(x, S(x)) \quad i = 1, \dots, o$$

Simulation Structure

$$f_i(x) = h_i(x, S(x)) \quad i = 1, \dots, o$$

Sum-of-squares structure:

$$h_i(x, S(x)) = \sum_{j \in N_i} (S_j(x))^2$$

where each N_1, \dots, N_o is an index set.

Increases order of approximation \Rightarrow
increases order of convergence

Simulation Structure

$$f_i(x) = h_i(x, S(x)) \quad i = 1, \dots, o$$

Sum-of-squares structure:

$$h_i(x, S(x)) = \sum_{j \in N_i} (S_j(x))^2$$

where each N_1, \dots, N_o is an index set.

Increases order of approximation \Rightarrow
increases order of convergence

Heterogeneous MOOPs:

$$\begin{aligned} h_1(x, S(x)) &= S_1(x) \\ h_2(x, S(x)) &= \|x\|^2 \end{aligned}$$

Use expensive surrogate models for h_1 (i.e.,
 S_1) but not for h_2

Flexible Problem Types

Flexible Problem Types

- ▶ Mixed variable-types

Flexible Problem Types

- ▶ Mixed variable-types **Ex: categorical variables**

Flexible Problem Types

- ▶ Mixed variable-types **Ex: categorical variables**
 - ▶ Focus on *embedding* into continuous space

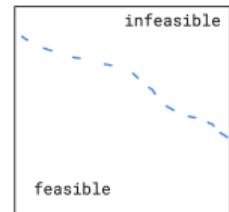
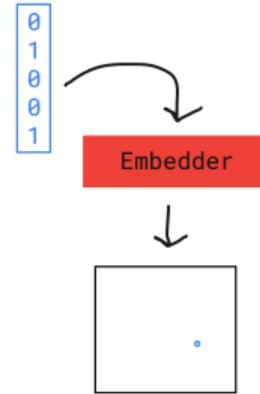
Flexible Problem Types

- ▶ Mixed variable-types **Ex: categorical variables**
 - ▶ Focus on *embedding* into continuous space
- ▶ (Nonlinear) constraints

Flexible Problem Types

- ▶ Mixed variable-types **Ex: categorical variables**
 - ▶ Focus on *embedding* into continuous space

- ▶ (Nonlinear) constraints
 - ▶ Focus on *augmented Lagrangian* penalties (relax to augmented unconstrained problem)



Design constraints

Design constraints

Easy to use:

- ▶ OOP gives a clean interface for adding attributes to a MOOP instance
- ▶ Total modularity (free to mix-and-match)

Design constraints

Easy to use:

- ▶ OOP gives a clean interface for adding attributes to a MOOP instance
- ▶ Total modularity (free to mix-and-match)

Easy to deploy:

- ▶ Simulation/experiment evaluations are only called in the `MOOP.solve()` method
- ▶ Extend `MOOP` class and overwrite `solve()` to deploy in different workflows
- ▶ **Ex:** Deploy parallel solvers on HPC systems using `libEnsemble`

Design constraints

Easy to use:

- ▶ OOP gives a clean interface for adding attributes to a MOOP instance
- ▶ Total modularity (free to mix-and-match)

Easy to deploy:

- ▶ Simulation/experiment evaluations are only called in the `MOOP.solve()` method
- ▶ Extend `MOOP` class and overwrite `solve()` to deploy in different workflows
- ▶ **Ex:** Deploy parallel solvers on HPC systems using `libEnsemble`

Easy to maintain and extend:

- ▶ OOP + total modularity makes adding new features easy
- ▶ Agile development with continuous integration
- ▶ Well-documented interface, contributing, and release process

ParMOO Release



Written in Python (available on pip and GitHub)



<https://parmo.readthedocs.io/en/latest/quickstart.html>



Combine with libEnsemble to use parallel solvers

Chang and Wild. ParMOO: A Python library for parallel multiobjective simulation optimization. *JOSS* 8(82):4468 (2023).

Example 1: Fayans EDF Model Calibration

Find params $x \in [0, 1]^{13}$ to fit the Fayans model to data d_i :

$$M(\xi_i; x) \approx d_i \quad i = 1, \dots, 198$$

ParMOO simulation:

$$S_i(x) = M(\xi_i; x) - d_i, \quad i = 1, \dots, 198;$$

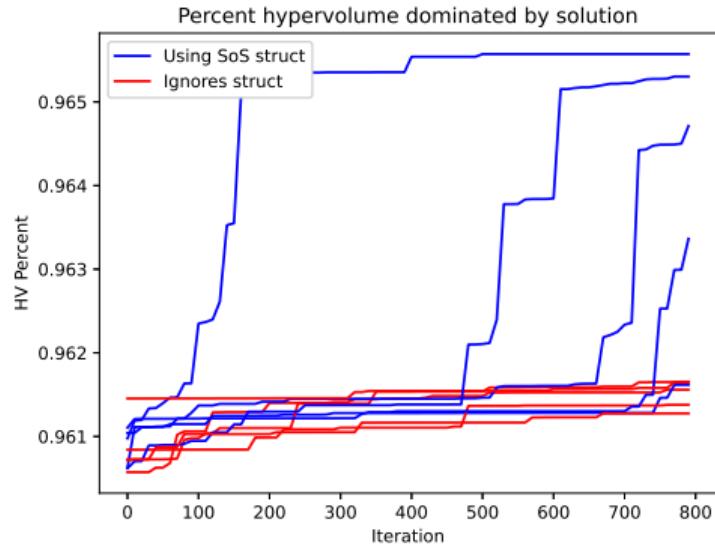
Min SOS across 3 observable classes

$$F_t = \sum_{i=1}^{m_t} (S_{t,i}(x))^2$$

Bollapragada et al. Journal of Physics G: Nuclear and Particle Physics 48(2), 2020.

Fayans Solution with ParMOO

- ▶ Approximated Fayans model using inv dist weighting on existing dataset
- ▶ Implemented parallel solver in ParMOO using libEnsemble
- ▶ Just **14-25 lines of Python code**
- ▶ Ran for **10K** sim evals
- ▶ Compared against **same solver w/o exploiting SOS structure**



Example 2: Material Manufacturing with ParMOO

Choose optimal settings for material manufacturing in a continuous flow reactor (CFR)

We know how to make a desired material, need to produce at scale:

1. **Maximize the product** (battery electrolyte: TFML)
2. Can increase temperature to **reduce reaction time**
3. Too much heat activates a side reaction; need to **minimize unwanted byproduct**

Challenges:

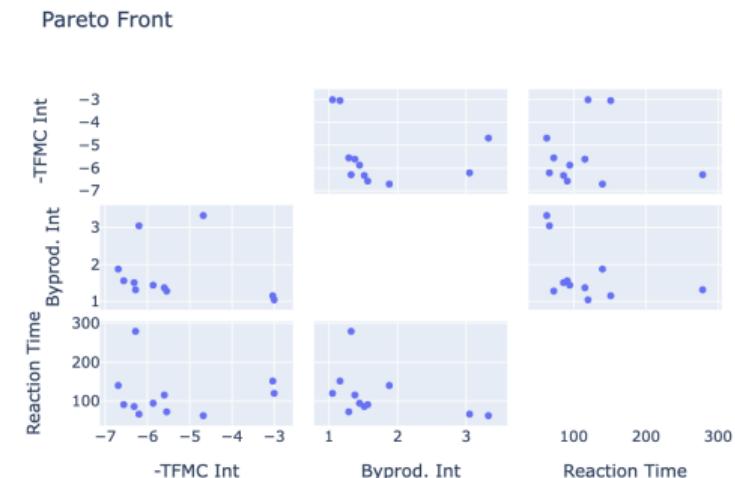
- ▶ Mixed variable types
- ▶ Heterogeneous objectives
- ▶ Must send experiments to run on CFR

CFR Optimization with ParMOO

Extend MOOP class to send/receive experiment data using MDML library
(Apache Kafka)

Used categorical variable embeddings

Modeled Product/Byproduct as simulations and reaction time using algebraic equation of input



Next Release

Coming in v. 0.2

- ▶ Interactive post-run visualization tools
- ▶ Support for customized embeddings and passing raw (unscaled) inputs
- ▶ MDML (Apache Kafka) interface for distributing simulation evaluations
- ▶ (Maybe) advanced techniques for design-of-experiments

Resources

E-mail: tchang@anl.gov
E-mail: parmoo@mcs.anl.gov

JOSS 8(82):4468 (2023)

GitHub: github.com/parmoo/parmoo
Docs: parmoo.readthedocs.io
PyPI: pip install parmoo

This material is based upon work supported by the U.S. Department of Energy, Office of Science, Office of Advanced Scientific Computing Research, SciDAC program under contract number DE-AC02-06CH11357.